

SVMs & K-NN

Shayling Zhao

1.**Support Vector Machines with Synthetic Data**

For this problem, we will generate synthetic data for a nonlinear binary classification problem and partition it into training, validation and test sets. Our goal is to understand the behavior of SVMs with Radial-Basis Function (RBF) kernels with different values of C and γ .

```
In [11]: import numpy as np
from sklearn.datasets import make_moons
from sklearn.model_selection import train_test_split
import matplotlib.pyplot as plt
from matplotlib.colors import ListedColormap
import os
from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsClassifier

def generate_data(n_samples, test_frac=0.2, val_frac=0.2):
    # Generate a Non-linear data set
    X, y = make_moons(n_samples=n_samples, noise=0.25, random_state=42)

    # Take a small subset of the data and make it VERY noisy; that is, generate outliers
    m = 30
    np.random.seed(30) # Deliberately use a different seed
    ind = np.random.permutation(n_samples)[:m]
    X[ind, :] += np.random.multivariate_normal([0, 0], np.eye(2), (m, ))
    y[ind] = 1 - y[ind]

    # Plot this data
    cmap = ListedColormap(['#300655', '#178000'])
    plt.scatter(X[:, 0], X[:, 1], c=y, cmap=cmap, edgecolors='k')

    # Plot it, we use train_test_split to partition (X, y) into training and test sets
    X_trn, X_tst, y_trn, y_tst = train_test_split(X, y, test_size=test_frac,
                                                random_state=42)

    # Next, we use train_test_split to further partition (X_trn, y_trn) into training and validation sets
    X_trn, X_val, y_trn, y_val = train_test_split(X_trn, y_trn, test_size=val_frac,
                                                random_state=42)

    return (X_trn, y_trn), (X_val, y_val), (X_tst, y_tst)
```

```
In [12]: def visualize(models, param, X, y):
    # Initialize plotting
    if len(models) % 3 == 0:
        nrows = len(models) // 3
    else:
        nrows = len(models) // 3 + 1

    fig, axes = plt.subplots(nrows=nrows, ncols=3, figsize=(15, 5.0 * nrows))
    cmap = ListedColormap(['#300655', '#178000'])

    # Create a mesh
    xMin, xMax = X[:, 0].min() - 1, X[:, 0].max() + 1
    yMin, yMax = X[:, 1].min() - 1, X[:, 1].max() + 1
    xMesh, yMesh = np.meshgrid(np.arange(xMin, xMax, 0.01),
                                np.arange(yMin, yMax, 0.01))

    for i, (p, clf) in enumerate(models.items()):
        # if i > 0:
        #     break
        r, c = np.divmod(i, 3)
        ax = axes[r, c]

        # Plot contours
        zMesh = clf.decision_function(np.c_[xMesh.ravel(), yMesh.ravel()])
        zMesh = zMesh.reshape(xMesh.shape)
        ax.contourf(xMesh, yMesh, zMesh, cmap=plt.cm.PYG, alpha=0.6)

        if (param == 'C' and p > 0.0) or (param == 'gamma'):
            ax.contour(xMesh, yMesh, zMesh, colors='k', levels=[-1, 0, 1],
                      alpha=0.5, linestyle=['--', '-', '--'])

    # Plot data
    ax.scatter(X[:, 0], X[:, 1], c=y, cmap=cmap, edgecolors='k')
    ax.set_title('{0} = {1}'.format(param, p))

In [13]: # Generate the data
n_samples = 300 # Total size of data set
(X_trn, y_trn), (X_val, y_val), (X_tst, y_tst) = generate_data(n_samples)
```

a. (25 points) The effect of the regularization parameter, C

Python code snippet below takes the generated synthetic 2-d data as input and learns non-linear SVMs. Used scikit-learn's *SVC* function to learn SVM models with **radial-basis kernels** for fixed γ and various choices of $C \in \{10^{-3}, 10^{-2}, \dots, 1, \dots, 10^3\}$. The value of γ is fixed to $\gamma = \frac{1}{4\sigma^2}$, where d is the data dimension and σ^2 is the standard deviation of the data set. X- SVC can automatically use these setting for γ if you pass the argument gamma = 'scale' (see documentation for more details).

```
In [14]: from sklearn.svm import SVC
# Learn support vector classifiers with a radial-basis function kernel with
# fixed gamma = 1 / (n_features * X.std()) and different values of C

def visualize(models, param, X, y):
    # Initialize plotting
    if len(models) % 3 == 0:
        nrows = len(models) // 3
    else:
        nrows = len(models) // 3 + 1

    fig, axes = plt.subplots(nrows=nrows, ncols=3, figsize=(15, 5.0 * nrows))
    cmap = ListedColormap(['#300655', '#178000'])

    # Create a mesh
    xMin, xMax = X[:, 0].min() - 1, X[:, 0].max() + 1
    yMin, yMax = X[:, 1].min() - 1, X[:, 1].max() + 1
    xMesh, yMesh = np.meshgrid(np.arange(xMin, xMax, 0.01),
                                np.arange(yMin, yMax, 0.01))

    for i, (p, clf) in enumerate(models.items()):
        # if i > 0:
        #     break
        r, c = np.divmod(i, 3)
        ax = axes[r, c]

        # Plot contours
        zMesh = clf.decision_function(np.c_[xMesh.ravel(), yMesh.ravel()])
        zMesh = zMesh.reshape(xMesh.shape)
        ax.contourf(xMesh, yMesh, zMesh, cmap=plt.cm.PYG, alpha=0.6)

        if (param == 'C' and p > 0.0) or (param == 'gamma'):
            ax.contour(xMesh, yMesh, zMesh, colors='k', levels=[-1, 0, 1],
                      alpha=0.5, linestyle=['--', '-', '--'])

    # Plot data
    ax.scatter(X[:, 0], X[:, 1], c=y, cmap=cmap, edgecolors='k')
    ax.set_title('{0} = {1}'.format(param, p))

C_range = np.arange(-3.0, 6.0, 1.0)
C_values = np.power(10.0, C_range)

models = dict()
trnErr = dict()
valErr = dict()

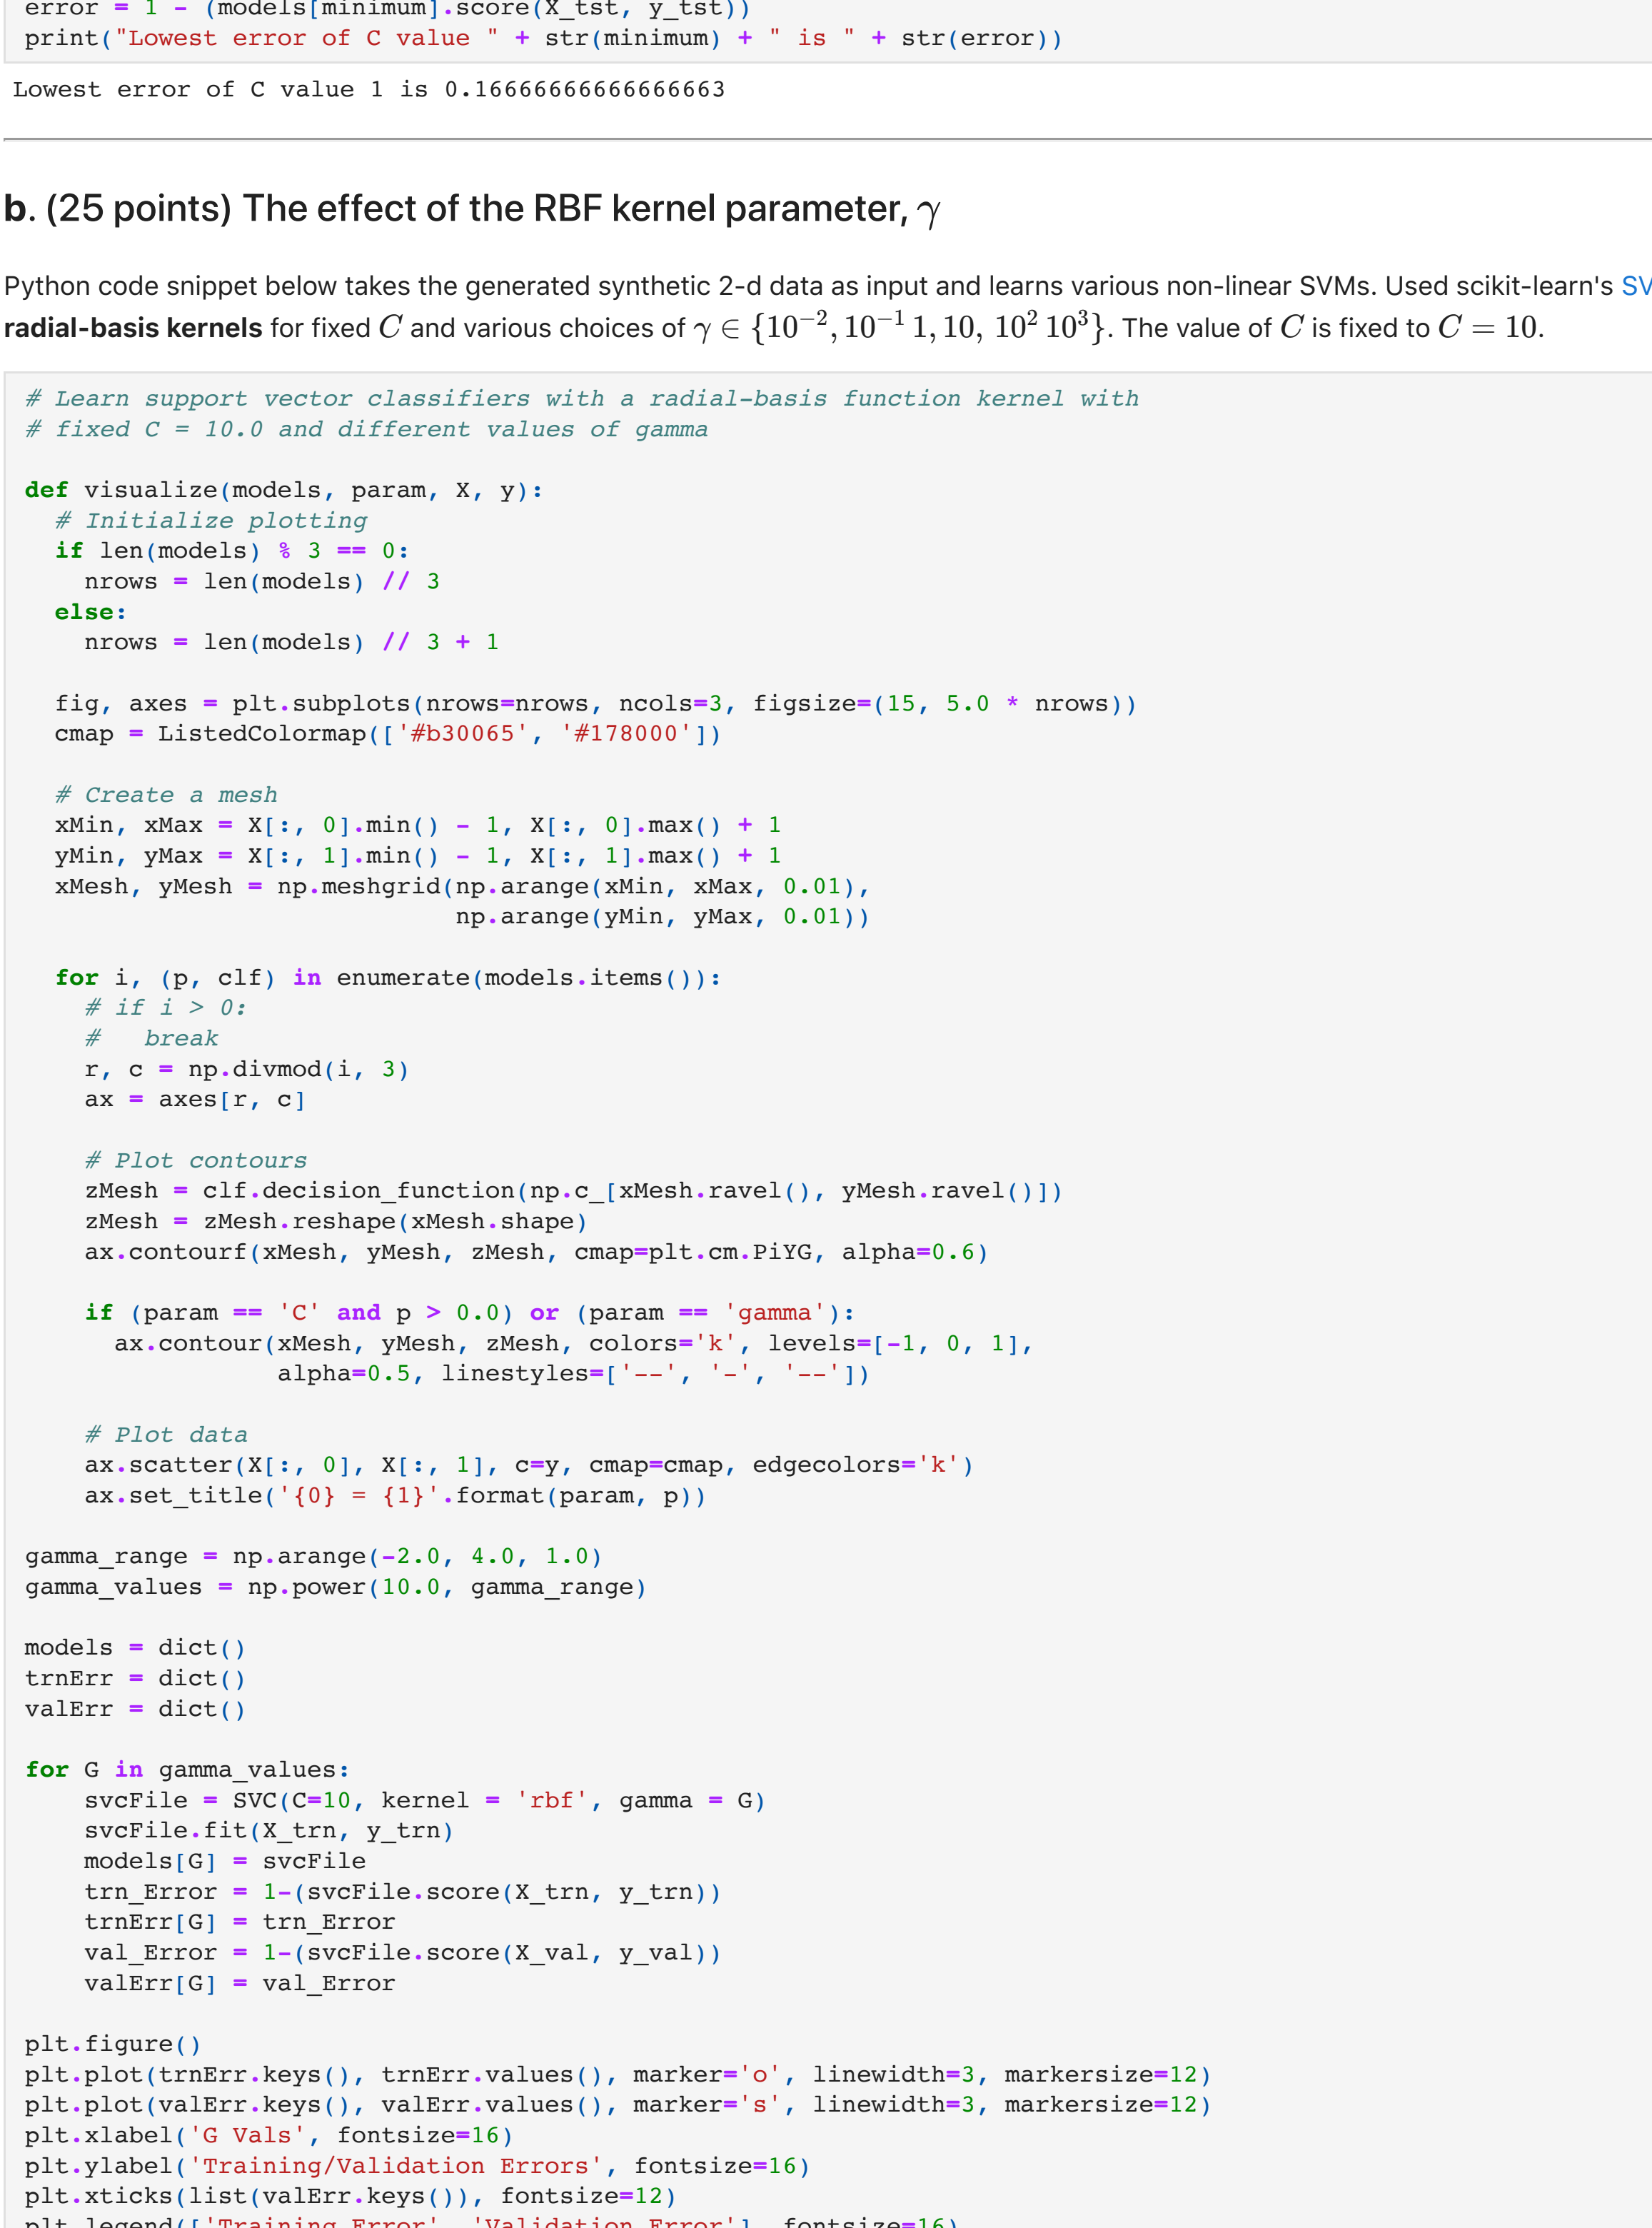
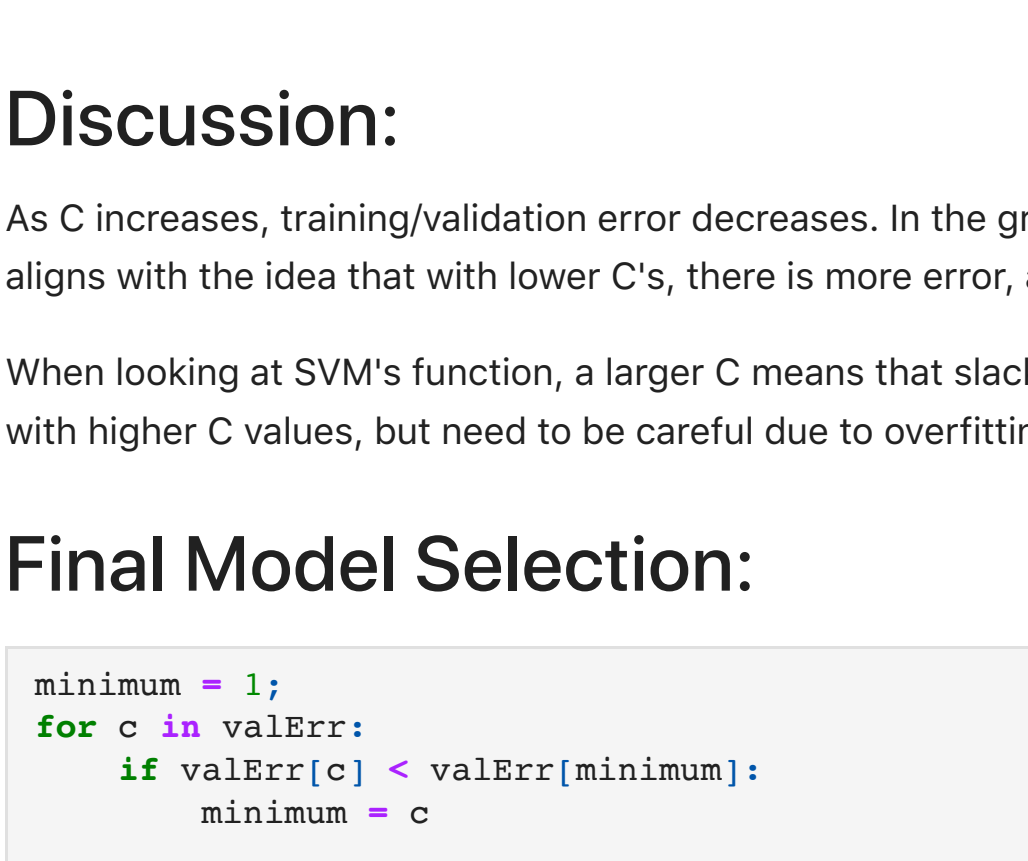
for C in C_values:
    svcFile = SVC(C=C, kernel='rbf', gamma='scale')
    svcFile.fit(X_trn, y_trn)
    models[C] = svcFile
    trn_Error = 1-(svcFile.score(X_trn, y_trn))
    trnErr[C] = trn_Error
    val_Error = 1-(svcFile.score(X_val, y_val))
    valErr[C] = val_Error

plt.figure()
plt.plot(trnErr.keys(), trnErr.values(), marker='o', linewidth=3, markersize=12)
plt.plot(valErr.keys(), valErr.values(), marker='s', linewidth=3, markersize=12)
plt.xlabel('C Vals', fontsize=16)
plt.ylabel('Training/Validation Errors', fontsize=16)
plt.xticks(list(valErr.keys()), fontsize=12)
plt.legend(['Training Error', 'Validation Error'], fontsize=16)
plt.xscale('log')

visualize(models, 'C', X_trn, y_trn)

minimum = 1;
for c in valErr:
    if models[c].score(X_tst, y_tst) > models[minimum].score(X_tst, y_tst):
        minimum = c

accuracy = models[minimum].score(X_tst, y_tst)
print("Accuracy " + str(minimum) + " is " + str(accuracy))
C Accuracy 100.0 is 0.85
```



Discussion:

As C increases, training/validation error decreases. In the graphs shown, you can see that as C increases, the complexity of shapes around the data points also increases. This aligns with the idea that with lower C 's, there is more error, and with higher C 's, there is less errors (but more likely to overfit).

When looking at SVM's function, a larger C means that slack variables becomes more penalized and a smaller margin accepted. Therefore, minimization of the function occurs with higher C values, but need to be careful due to overfitting.

Final Model Selection:

```
In [15]: minimum = 1;
for c in valErr:
    if valErr[c] < valErr[minimum]:
        minimum = c

error = 1 - (models[minimum].score(X_tst, y_tst))
print("Lowest error of C value " + str(minimum) + " is " + str(error))

Lowest error of C value 1 is 0.16666666666666663
```

b. (25 points) The effect of the RBF kernel parameter, γ

Python code snippet below takes the generated synthetic 2-d data as input and learns various non-linear SVMs. Used scikit-learn's *SVC* function to learn SVM models with **radial-basis kernels** for fixed C and various choices of $\gamma \in \{10^{-2}, 10^{-1}, 1, 10, 10^2, 10^3\}$. The value of C is fixed to $C = 10$.

```
In [16]: # Learn support vector classifiers with a radial-basis function kernel with
# fixed C = 10.0 and different values of gamma

def visualize(models, param, X, y):
    # Initialize plotting
    if len(models) % 3 == 0:
        nrows = len(models) // 3
    else:
        nrows = len(models) // 3 + 1

    fig, axes = plt.subplots(nrows=nrows, ncols=3, figsize=(15, 5.0 * nrows))
    cmap = ListedColormap(['#300655', '#178000'])

    # Create a mesh
    xMin, xMax = X[:, 0].min() - 1, X[:, 0].max() + 1
    yMin, yMax = X[:, 1].min() - 1, X[:, 1].max() + 1
    xMesh, yMesh = np.meshgrid(np.arange(xMin, xMax, 0.01),
                                np.arange(yMin, yMax, 0.01))

    for i, (p, clf) in enumerate(models.items()):
        # if i > 0:
        #     break
        r, c = np.divmod(i, 3)
        ax = axes[r, c]

        # Plot contours
        zMesh = clf.decision_function(np.c_[xMesh.ravel(), yMesh.ravel()])
        zMesh = zMesh.reshape(xMesh.shape)
        ax.contourf(xMesh, yMesh, zMesh, cmap=plt.cm.PYG, alpha=0.6)

        if (param == 'C' and p > 0.0) or (param == 'gamma'):
            ax.contour(xMesh, yMesh, zMesh, colors='k', levels=[-1, 0, 1],
                      alpha=0.5, linestyle=['--', '-', '--'])

    # Plot data
    ax.scatter(X[:, 0], X[:, 1], c=y, cmap=cmap, edgecolors='k')
    ax.set_title('{0} = {1}'.format(param, p))

gamma_range = np.arange(-2.0, 4.0, 1.0)
gamma_values = np.power(10.0, gamma_range)

models = dict()
trnErr = dict()
valErr = dict()

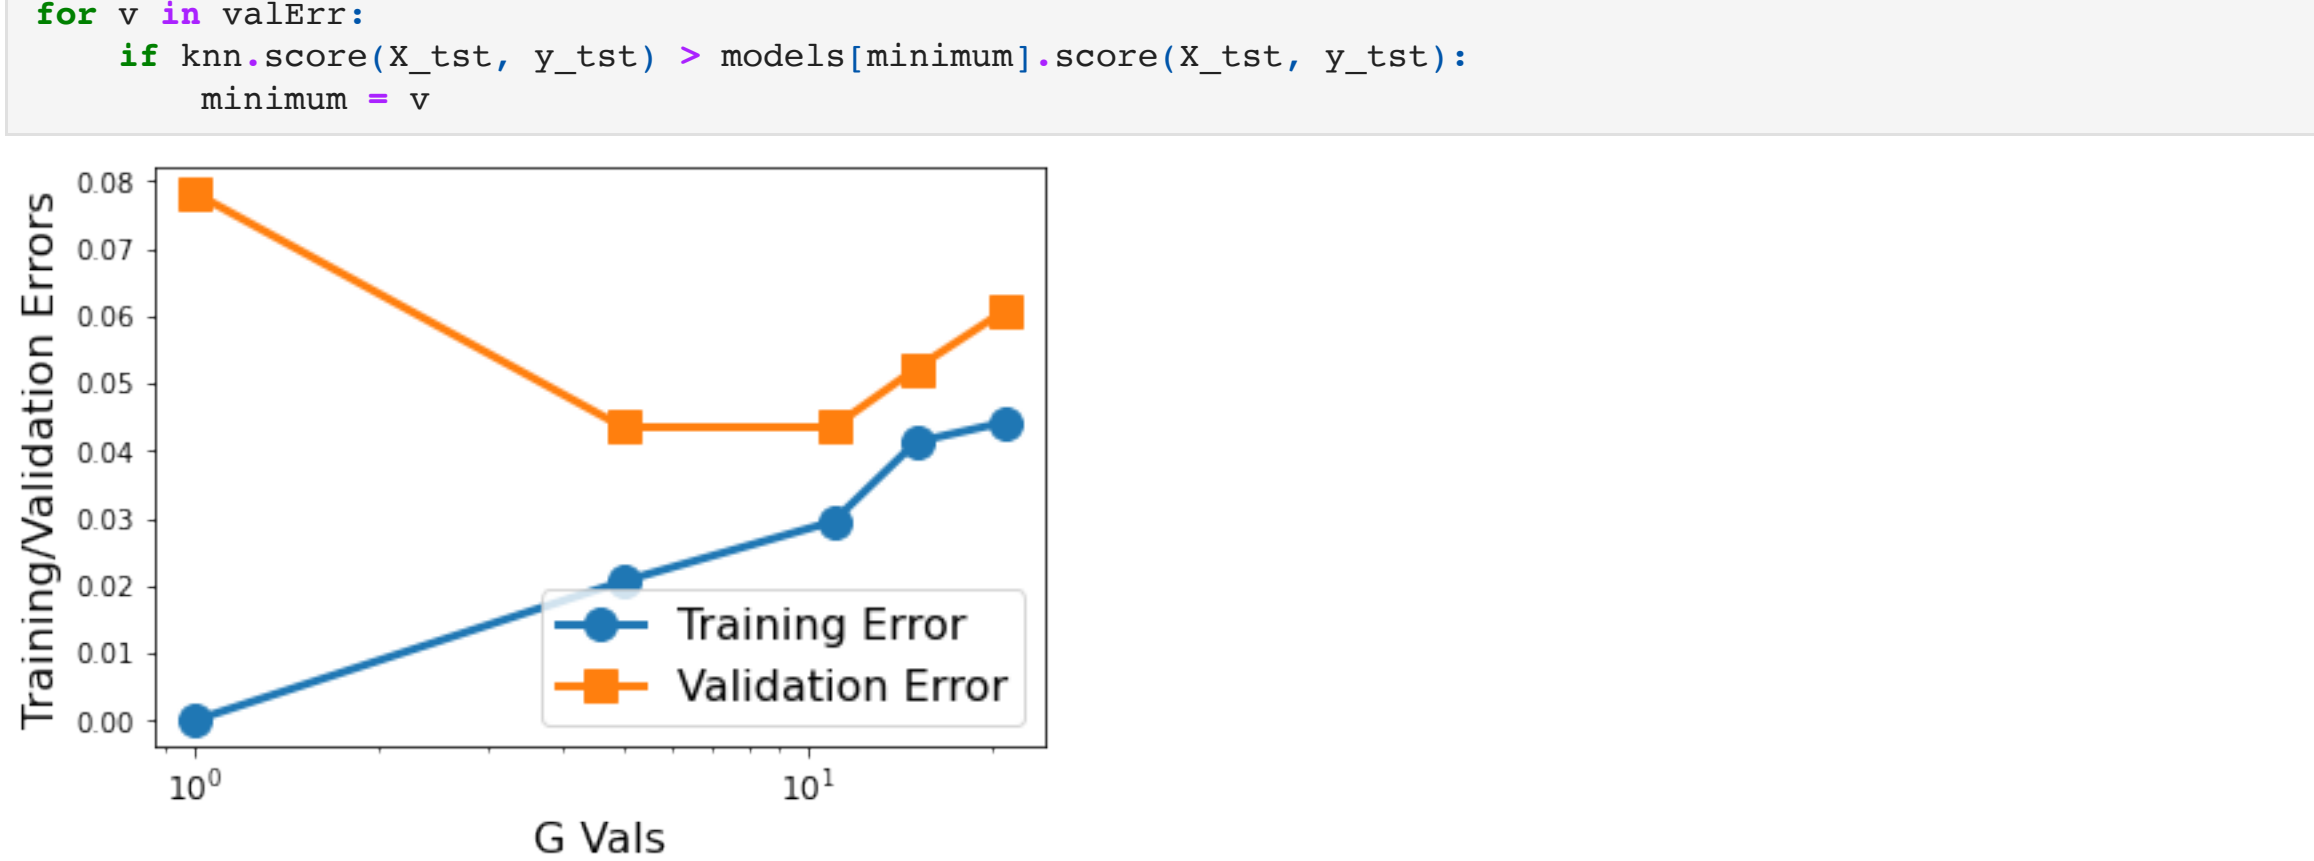
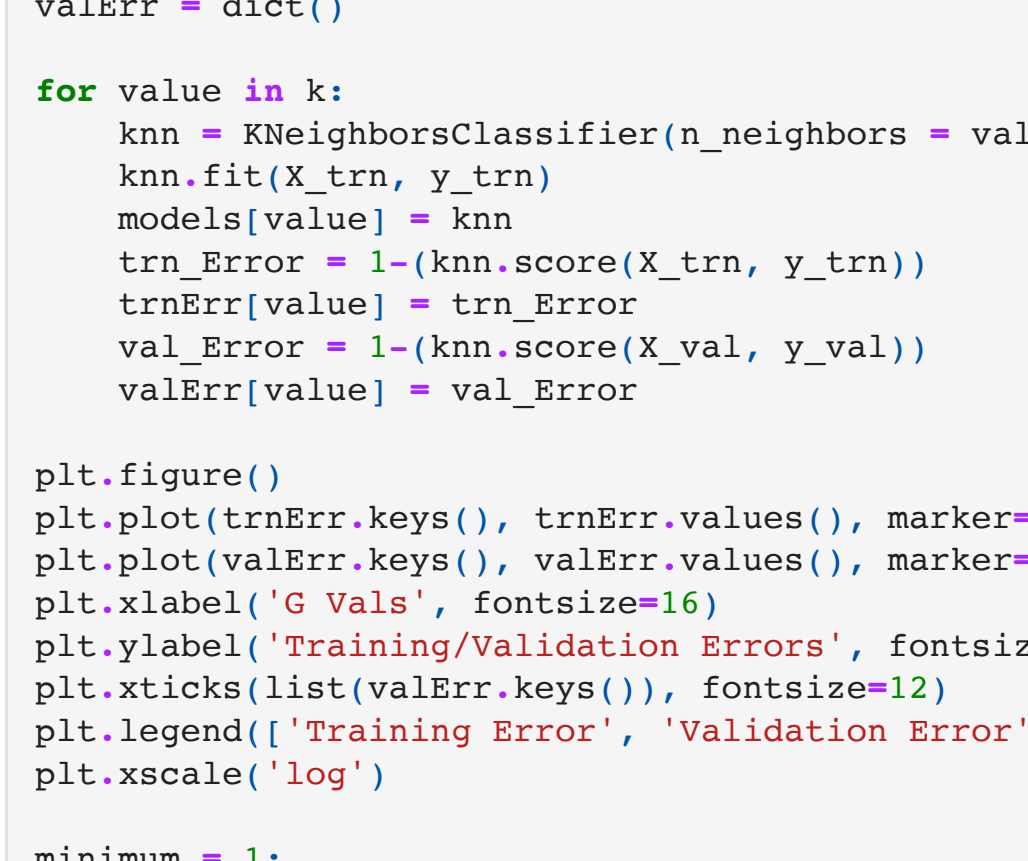
for G in gamma_values:
    svcFile = SVC(C=10, kernel='rbf', gamma=G)
    svcFile.fit(X_trn, y_trn)
    models[G] = svcFile
    trn_Error = 1-(svcFile.score(X_trn, y_trn))
    trnErr[G] = trn_Error
    val_Error = 1-(svcFile.score(X_val, y_val))
    valErr[G] = val_Error

plt.figure()
plt.plot(trnErr.keys(), trnErr.values(), marker='o', linewidth=3, markersize=12)
plt.plot(valErr.keys(), valErr.values(), marker='s', linewidth=3, markersize=12)
plt.xlabel('G Vals', fontsize=16)
plt.ylabel('Training/Validation Errors', fontsize=16)
plt.xticks(list(valErr.keys()), fontsize=12)
plt.legend(['Training Error', 'Validation Error'], fontsize=16)
plt.xscale('log')

visualize(models, 'gamma', X_trn, y_trn)

minimum = 1;
for g in valErr:
    if models[g].score(X_tst, y_tst) > models[minimum].score(X_tst, y_tst):
        minimum = g

accuracy = models[minimum].score(X_tst, y_tst)
print("Accuracy " + str(minimum) + " is " + str(accuracy))
G Accuracy 1 is 0.8333333333333334
```



Discussion:

As gamma increases, the training error decreases. As gamma decreases, the validation error decreases but increases as gamma values increase. In SVC, a consequence of higher gamma values is that the model tries to fit the training set.

As seen in the models, the area around points/the data set are more constricted/specific as values of gamma increases. In the training/validation error chart, the validation error starts off in a decrease, but changes course to increase after the gamma values becomes 10 or more. Meanwhile, the training error continues on a consistent decline as the gamma values increase.

In relation to the functional form of the RBF kernel, small gamma yields a broad decision region while large gamma yields a more constricted/specific decision region because the higher the gamma value, the equation will yield a lower $K(x,z)$.

Final Model Selection:

```
In [17]: minimum = 1;
for g in valErr:
    if valErr[g] < valErr[minimum]:
        minimum = g

error = 1 - (models[minimum].score(X_tst, y_tst))
print("Lowest error of G value " + str(minimum) + " is " + str(1-error))
print("So, the highest test accuracy is " + str(error))

Lowest error of G value (100.0, 0.01) is 0.03478268695652195
So, the highest test accuracy is 0.9652173913043478
```

2.**Breast Cancer Diagnosis with Support Vector Machines**, 25 points.

For this problem, we will use the *Wisconsin Breast Cancer* data set, which has already been pre-processed and partitioned into training, validation and test sets. Numpy's *loadtxt* command can be used to load CSV files.

```
In [18]: train = np.loadtxt('C:/Users/shaylinghao/Desktop/wdbc_trn.csv', delimiter=",")
X_trn = train[:,1:]
y_trn = train[:,0]
test = np.loadtxt('C:/Users/shaylinghao/Desktop/wdbc_val.csv', delimiter=",")
X_tst = test[:,1:]
y_tst = test[:,0]
val = np.loadtxt('C:/Users/shaylinghao/Desktop/wdbc_val.csv', delimiter=",")
X_val = val[:,1:]
y_val = val[:,0]

# Used scikit-learn's SVC function to learn SVM models with radial-basis kernels for each combination of C in {10^-2, 10^-1, 1, 10^1, ... 10^3} and
gamma in {10^-2, 10^-1, 1, 10, 10^2}.
```

```
In [19]: C_range = np.arange(-3.0, 6.0, 1.0)
C_values = np.power(10.0, C_range)

gamma_range = np.arange(-2.0, 4.0, 1.0)
gamma_values = np.power(10.0, gamma_range)

models = dict()
trnErr = dict()
valErr = dict()

for C in C_values:
    for G in gamma_values:
        svcFile = SVC(C=C, kernel='rbf', gamma=G)
        svcFile.fit(X_trn, y_trn)
        models[(C, G)] = svcFile
        trn_Error = 1-(svcFile.score(X_trn, y_trn))
        trnErr[(C, G)] = trn_Error
        val_Error = 1-(svcFile.score(X_val, y_val))
        valErr[(C, G)] = val_Error

trnErr = {}
cValErr = {}
for C in C_values:
    trn_sum = 0
    val_sum = 0
    for G in gamma_values:
        trn_sum += trnErr[(C, G)]
        val_sum += valErr[(C, G)]
    cValErr[C] = trn_sum/len(C_values)
    cValErr[C] = val_sum/len(gamma_values)

plt.figure()
plt.plot(trnErr.keys(), trnErr.values(), marker='o', linewidth=3, markersize=12)
plt.plot(cValErr.keys(), cValErr.values(), marker='s', linewidth=3, markersize=12)
plt.xlabel('G Vals', fontsize=16)
plt.ylabel('Training/Validation Errors', fontsize=16)
plt.xticks(list(cValErr.keys()), fontsize=12)
plt.legend(['Training Error', 'Validation Error'], fontsize=16)
plt.xscale('log')

gTrnErr = {}
gValErr = {}
for G in gamma_values:
    gtrn_sum = 0
    gval_sum = 0
    for C in C_values:
        gtrn_sum += trnErr[(C, G)]
        gval_sum += valErr[(C, G)]
    gTrnErr[G] = gtrn_sum/len(C_values)
    gValErr[G] = gval_sum/len(C_values)

plt.figure()
plt.plot(gTrnErr.keys(), gTrnErr.values(), marker='o', linewidth=3, markersize=12)
plt.plot(gValErr.keys(), gValErr.values(), marker='s', linewidth=3, markersize=12)
plt.xlabel('C Vals', fontsize=16)
plt.ylabel('Training/Validation Errors', fontsize=16)
plt.xticks(list(gValErr.keys()), fontsize=12)
plt.legend(['Training Error', 'Validation Error'], fontsize=16)
plt.xscale('log')

minimum = list(valErr.keys())[0]
for key in valErr.keys():
    if models[key].score(X_tst, y_tst) > models[minimum].score(X_tst, y_tst):
        minimum = key

accuracy = models[minimum].score(X_tst, y_tst)
print("G and C Accuracy " + str(minimum) + " is " + str(accuracy))

G and C Accuracy (10.0, 0.01) is 0.9826086956521739
```


Final Model Selection:

```
In [22]: accuracy = models[minimum].score(X_tst, y_tst)
print("Best K value = " + str(minimum) + " and has an accuracy of " + str(accuracy))

Best K value = 11 and has an accuracy of 0.9739130434782609
```

Discussion:

Discussion: Which of these two approaches, SVMs or kNN, would you prefer for this classification task? Explain.

The kNN classification task is preferred because it yielded a higher accuracy percentage of 97.39% compared to the SVM that yielded an accuracy percentage of 95.65%. kNN also works better with data that has many attributes while SVM needs to match C with gamma which could require additional work/time.